

In [1]:

```

1 import itertools
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from matplotlib.ticker import NullFormatter
5 import pandas as pd
6 import numpy as np
7 import matplotlib.ticker as ticker
8 from sklearn import preprocessing
9 %matplotlib inline

```

In [3]:

```

1 df = pd.read_csv('loan_train.csv')
2 df.head()

```

Out[3]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	0	0	PAIDOFF	1000	30	9/8/2016	10/7/2016	45	High School or Below
1	2	2	PAIDOFF	1000	30	9/8/2016	10/7/2016	33	Bechal
2	3	3	PAIDOFF	1000	15	9/8/2016	9/22/2016	27	college
3	4	4	PAIDOFF	1000	30	9/9/2016	10/8/2016	28	college
4	6	6	PAIDOFF	1000	30	9/9/2016	10/8/2016	29	college

In [4]:

```
1 df.shape
```

Out[4]:

(346, 10)

Convert to date time object

In [5]:

```

1 df['due_date'] = pd.to_datetime(df['due_date'])
2 df['effective_date'] = pd.to_datetime(df['effective_date'])
3 df.head()

```

Out[5]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	High School or Below
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bechalo
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	college
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	college
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	college

In [8]:

```
1 df.dtypes
```

Out[8]:

```

Unnamed: 0          int64
Unnamed: 0.1        int64
loan_status        object
Principal          int64
terms              int64
effective_date      datetime64[ns]
due_date            datetime64[ns]
age                int64
education           object
Gender             object
dtype: object

```

Data visualization and pre-processing

Let's see how many of each class is in our data set

In [9]:

```
1 df['loan_status'].value_counts()
```

Out[9]:


```
PAIDOFF      260
COLLECTION    86
Name: loan_status, dtype: int64
```

260 people have paid off the loan on time while 86 have gone into collection

Lets plot some columns to understand data better:


In [10]:

```
1 import seaborn as sns
2
3 bins = np.linspace(df.Principal.min(), df.Principal.max(), 10)
4 g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
5 g.map(plt.hist, 'Principal', bins=bins, ec="k")
6
7 g.axes[-1].legend()
8 plt.show()
```



In [11]:

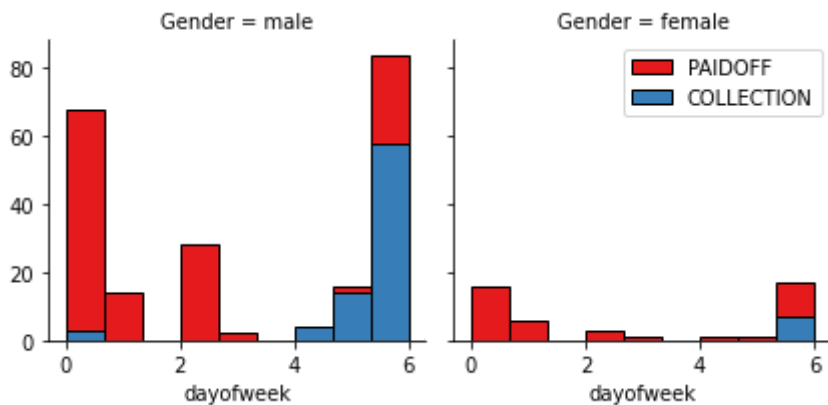
```
1 bins = np.linspace(df.age.min(), df.age.max(), 10)
2 g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
3 g.map(plt.hist, 'age', bins=bins, ec="k")
4
5 g.axes[-1].legend()
6 plt.show()
```



Lets look at the day of the week people get the loan

In [12]:

```
1 df['dayofweek'] = df['effective_date'].dt.dayofweek
2 bins = np.linspace(df.dayofweek.min(), df.dayofweek.max(), 10)
3 g = sns.FacetGrid(df, col="Gender", hue="loan_status", palette="Set1", col_wrap=2)
4 g.map(plt.hist, 'dayofweek', bins=bins, ec="k")
5 g.axes[-1].legend()
6 plt.show()
7
```



We see that people who get the loan at the end of the week dont pay it off, so lets use Feature binarization to set a threshold values less then day 4

In [13]:

```
1 df['weekend'] = df['dayofweek'].apply(lambda x: 1 if (x>3) else 0)
2 df.head()
```

Out[13]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	High School or Below
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bechalo
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	college
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	college
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	college

Convert Categorical features to numerical values

In [14]:

```
1 df.groupby(['Gender'])['loan_status'].value_counts(normalize=True)
```

Out[14]:

```
Gender  loan_status
female  PAIDOFF      0.865385
        COLLECTION  0.134615
male    PAIDOFF      0.731293
        COLLECTION  0.268707
Name: loan_status, dtype: float64
```

86 % of female pay there loans while only 73 % of males pay there loan

In [17]:

```
1 df['Gender'].replace({'Female':1,'Male':0},inplace=True)
2 df
```

Out[17]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	educ
0	0	0	PAIDOFF	1000	30	2016-09-08	2016-10-07	45	SchE
1	2	2	PAIDOFF	1000	30	2016-09-08	2016-10-07	33	Bec
2	3	3	PAIDOFF	1000	15	2016-09-08	2016-09-22	27	cc
3	4	4	PAIDOFF	1000	30	2016-09-09	2016-10-08	28	cc
4	6	6	PAIDOFF	1000	30	2016-09-09	2016-10-08	29	cc
...
341	394	394	COLLECTION	800	15	2016-09-11	2016-09-25	32	SchE
342	395	395	COLLECTION	1000	30	2016-09-11	2016-10-10	25	SchE
343	397	397	COLLECTION	800	15	2016-09-12	2016-09-26	39	cc
344	398	398	COLLECTION	1000	30	2016-09-12	2016-11-10	28	cc
345	399	399	COLLECTION	1000	30	2016-09-12	2016-10-11	26	cc

346 rows × 12 columns



One Hot Encoding How about education?

In [18]:

```
1 df.groupby(['education'])['loan_status'].value_counts(normalize=True)
```

Out[18]:

education	loan_status	
Bechalor	PAIDOFF	0.750000
	COLLECTION	0.250000
High School or Below	PAIDOFF	0.741722
	COLLECTION	0.258278
Master or Above	COLLECTION	0.500000
	PAIDOFF	0.500000
college	PAIDOFF	0.765101
	COLLECTION	0.234899

Name: loan_status, dtype: float64

Use one hot encoding technique to conver categorical variables to binary variables and append them to the feature Data Frame

In [19]:

```
1 Feature = df[['Principal','terms','age','Gender','weekend']]
2 Feature = pd.concat([Feature,pd.get_dummies(df['education'])], axis=1)
3 Feature.drop(['Master or Above'], axis = 1,inplace=True)
4 Feature.head()
```

Out[19]:

	Principal	terms	age	Gender	weekend	Bechalor	High School or Below	college
0	1000	30	45	0	0	0	1	0
1	1000	30	33	1	0	1	0	0
2	1000	15	27	0	0	0	0	1
3	1000	30	28	1	1	0	0	1
4	1000	30	29	0	1	0	0	1

In [20]:

```
1 X = Feature
2 X[0:5]
```

Out[20]:

	Principal	terms	age	Gender	weekend	Bechalor	High School or Below	college
0	1000	30	45	0	0	0	1	0
1	1000	30	33	1	0	1	0	0
2	1000	15	27	0	0	0	0	1
3	1000	30	28	1	1	0	0	1
4	1000	30	29	0	1	0	0	1

In [21]:

```
1 y = df['loan_status'].values
2 y[0:5]
```

Out[21]:

```
array(['PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF', 'PAIDOFF'],
      dtype=object)
```

Normalize Data

In [22]:

```
1 X= preprocessing.StandardScaler().fit(X).transform(X)
2 X[0:5]
```

Out[22]:

```
array([[ 0.51578458,  0.92071769,  2.33152555, -0.42056004, -1.20577805,
        -0.38170062,  1.13639374, -0.86968108],
       [ 0.51578458,  0.92071769,  0.34170148,  2.37778177, -1.20577805,
         2.61985426, -0.87997669, -0.86968108],
       [ 0.51578458, -0.95911111, -0.65321055, -0.42056004, -1.20577805,
        -0.38170062, -0.87997669,  1.14984679],
       [ 0.51578458,  0.92071769, -0.48739188,  2.37778177,  0.82934003,
        -0.38170062, -0.87997669,  1.14984679],
       [ 0.51578458,  0.92071769, -0.3215732 , -0.42056004,  0.82934003,
        -0.38170062, -0.87997669,  1.14984679]])
```

Classification

Now, it is your turn, use the training set to build an accurate model. Then use the test set to report the accuracy of the model You should use the following algorithm:

K Nearest Neighbor(KNN) Decision Tree Support Vector Machine Logistic Regression __ Notice: __

You can go above and change the pre-processing, feature selection, feature-extraction, and so on, to make a better model. You should use either scikit-learn, Scipy or Numpy libraries for developing the classification algorithms. You should include the code of the algorithm in the following cells.

Train Test split

In [23]:

```
1 from sklearn.model_selection import train_test_split
2 x_train, x_test, y_train, y_test = train_test_split( X, y, test_size=0.2, random_state=
3 print ('Train set:', x_train.shape, y_train.shape)
4 print ('Test set:', x_test.shape, y_test.shape)
```

Train set: (276, 8) (276,)

Test set: (70, 8) (70,)

K Nearest Neighbor(KNN)

Notice: You should find the best k to build the model with the best accuracy. warning: You should not use the loan_test.csv for finding the best k, however, you can split your train_loan.csv into train and test to find the best k.

In [24]:

```
1 from sklearn.neighbors import KNeighborsClassifier
2 from sklearn.metrics import accuracy_score
```

Checking for the best value of K

In [25]:

```
1 for k in range(1, 10):
2     knn_model = KNeighborsClassifier(n_neighbors = k).fit(x_train, y_train)
3     knn_yhat = knn_model.predict(x_test)
4     print("For K = {} accuracy = {}".format(k, accuracy_score(y_test, knn_yhat)))
```

```
For K = 1 accuracy = 0.6714285714285714
For K = 2 accuracy = 0.6571428571428571
For K = 3 accuracy = 0.7142857142857143
For K = 4 accuracy = 0.6857142857142857
For K = 5 accuracy = 0.7571428571428571
For K = 6 accuracy = 0.7142857142857143
For K = 7 accuracy = 0.7857142857142857
For K = 8 accuracy = 0.7571428571428571
For K = 9 accuracy = 0.7571428571428571
```

In [26]:

```
1 print("We can see that the KNN model is the best for K=7")
```

We can see that the KNN model is the best for K=7

Building the model with the best value of K = 7

In [27]:

```
1 best_knn_model = KNeighborsClassifier(n_neighbors = 7).fit(x_train, y_train)
2 best_knn_model
```

Out[27]:

KNeighborsClassifier(n_neighbors=7)

In [29]:

```
1 ## Evaluation Metric
2
3 from sklearn.metrics import f1_score
4
5 print("Train set Accuracy (F1): ", f1_score(y_train, best_knn_model.predict(x_train), average='micro'))
6 print("Test set Accuracy (F1): ", f1_score(y_test, best_knn_model.predict(x_test), average='micro'))
```

```
Train set Accuracy (F1): 0.8000194668761034
Test set Accuracy (F1): 0.7766540244416351
```


Decision Tree

In [30]:

```
1 # importing libraries
2 from sklearn.tree import DecisionTreeClassifier
```

In [31]:

```
1 for d in range(1,10):
2     dt = DecisionTreeClassifier(criterion = 'entropy', max_depth = d).fit(x_train, y_train)
3     dt_yhat = dt.predict(x_test)
4     print("For depth = {} the accuracy score is {}".format(d, accuracy_score(y_test, dt_yhat)))
```

```
For depth = 1 the accuracy score is 0.7857142857142857
For depth = 2 the accuracy score is 0.7857142857142857
For depth = 3 the accuracy score is 0.6142857142857143
For depth = 4 the accuracy score is 0.6142857142857143
For depth = 5 the accuracy score is 0.6428571428571429
For depth = 6 the accuracy score is 0.7714285714285715
For depth = 7 the accuracy score is 0.7571428571428571
For depth = 8 the accuracy score is 0.7571428571428571
For depth = 9 the accuracy score is 0.6571428571428571
```

In [32]:

```
1 print("The best value of depth is d = 2 ")
```

The best value of depth is d = 2

In [33]:

```
1 ## Creating the best model for decision tree with best value of depth 2
2
3 best_dt_model = DecisionTreeClassifier(criterion = 'entropy', max_depth = 2).fit(x_train, y_train)
4 best_dt_model
```

Out[33]:

DecisionTreeClassifier(criterion='entropy', max_depth=2)

In [36]:

```
1
2 from sklearn.metrics import f1_score
3
4
5 print("Train set Accuracy (F1): ", f1_score(y_train, best_dt_model.predict(x_train), average='micro'))
6 print("Test set Accuracy (F1): ", f1_score(y_test, best_dt_model.predict(x_test), average='micro'))
```

```
Train set Accuracy (F1): 0.6331163939859591
Test set Accuracy (F1): 0.6914285714285714
```

Support Vector Machine

In [37]:

```
1 #importing svm
2 from sklearn import svm
3 from sklearn.metrics import f1_score
```

In [38]:

```
1 for k in ('linear', 'poly', 'rbf', 'sigmoid'):
2     svm_model = svm.SVC( kernel = k).fit(x_train,y_train)
3     svm_yhat = svm_model.predict(x_test)
4     print("For kernel: {}, the f1 score is: {}".format(k,f1_score(y_test,svm_yhat, aver
```

For kernel: linear, the f1 score is: 0.6914285714285714
For kernel: poly, the f1 score is: 0.7064793130366899
For kernel: rbf, the f1 score is: 0.7275882012724117
For kernel: sigmoid, the f1 score is: 0.6892857142857144

In [39]:

```
1 print("We can see the rbf has the best f1 score ")
```

We can see the rbf has the best f1 score

In [40]:

```
1 ## building best SVM with kernel = rbf
2 best_svm = svm.SVC(kernel='rbf').fit(x_train,y_train)
3 best_svm
```

Out[40]:

SVC()

In [42]:

```
1 ## Evaluation Metrics
2 # jaccard score and f1 score
3 from sklearn.metrics import f1_score
4
5
6 print("Train set Accuracy (F1): ", f1_score(y_train, best_svm.predict(x_train), average='v
7 print("Test set Accuracy (F1): ", f1_score(y_test, best_svm.predict(x_test), average='v
```

Train set Accuracy (F1): 0.7682165861513688
Test set Accuracy (F1): 0.7275882012724117

Logistic Regression

In [43]:

```
1 # importing libraries
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.metrics import log_loss
```

In [44]:

```

1 for k in ('lbfgs', 'saga', 'liblinear', 'newton-cg', 'sag'):
2     lr_model = LogisticRegression(C = 0.01, solver = k).fit(x_train, y_train)
3     lr_yhat = lr_model.predict(x_test)
4     y_prob = lr_model.predict_proba(x_test)
5     print('When Solver is {}, logloss is : {}'.format(k, log_loss(y_test, y_prob)))

```

When Solver is lbfgs, logloss is : 0.4920179847937498
 When Solver is saga, logloss is : 0.49201758723624645
 When Solver is liblinear, logloss is : 0.5772287609479654
 When Solver is newton-cg, logloss is : 0.4920178014679269
 When Solver is sag, logloss is : 0.492007407853154

In [45]:

```

1 print("We can see that the best solver is liblinear")

```

We can see that the best solver is liblinear

In [46]:

```

1 # Best Logistic regression model with Liblinear solver
2
3 best_lr_model = LogisticRegression(C = 0.01, solver = 'liblinear').fit(x_train, y_train)
4 best_lr_model

```

Out[46]:

LogisticRegression(C=0.01, solver='liblinear')

In [47]:

```

1 ## Evaluation Metrics
2 # jaccard score and f1 score
3 from sklearn.metrics import f1_score
4
5 print("Train set Accuracy (F1): ", f1_score(y_train, best_lr_model.predict(x_train), average='micro'))
6 print("Test set Accuracy (F1): ", f1_score(y_test, best_lr_model.predict(x_test), average='micro'))

```

Train set Accuracy (F1): 0.7341146337750953
 Test set Accuracy (F1): 0.6670522459996144

Model Evaluation using Test set

In [50]:

```

1
2 from sklearn.metrics import f1_score
3 from sklearn.metrics import log_loss

```

Load Test set for evaluation

In [52]:

```
1 test_df = pd.read_csv('loan_test.csv')
2 test_df.head()
```

Out[52]:

	Unnamed: 0	Unnamed: 0.1	loan_status	Principal	terms	effective_date	due_date	age	education
0	1	1	PAIDOFF	1000	30	9/8/2016	10/7/2016	50	Bechal
1	5	5	PAIDOFF	300	7	9/9/2016	9/15/2016	35	Master o Above
2	21	21	PAIDOFF	1000	30	9/10/2016	10/9/2016	43	Higl School o Belov
3	24	24	PAIDOFF	1000	30	9/10/2016	10/9/2016	26	colleg
4	35	35	PAIDOFF	800	15	9/11/2016	9/25/2016	29	Bechal

In [63]:

```
1 # data processing
2 test_df['due_date'] = pd.to_datetime(test_df['due_date'])
3 test_df['effective_date'] = pd.to_datetime(test_df['effective_date'])
4 test_df['dayofweek'] = test_df['effective_date'].dt.dayofweek
5
6 test_df['weekend'] = test_df['dayofweek'].apply(lambda x: 1 if (x>3) else 0)
7 test_df['Gender'].replace(to_replace=['male','female'], value=[0,1],inplace=True)
8
9 Feature1 = test_df[['Principal','terms','age','Gender','weekend']]
10 Feature1 = pd.concat([Feature1,pd.get_dummies(test_df['education'])], axis=1)
11 Feature1.drop(['Master or Above'], axis = 1,inplace=True)
12
13
14 x_loan_test = Feature1
15 x_loan_test = preprocessing.StandardScaler().fit(x_loan_test).transform(x_loan_test)
16
17 y_loan_test = test_df['loan_status'].values
```

In [64]:

```
1 from sklearn.metrics import f1_score
```

In [69]:

```
1 # F1_score
2
3 # KNN
4 knn_yhat = best_knn_model.predict(x_loan_test)
5 f1 = round(f1_score(y_loan_test, knn_yhat, average = 'weighted'), 2)
6
7 # Decision Tree
8 dt_yhat = best_dt_model.predict(x_loan_test)
9 f2 = round(f1_score(y_loan_test, dt_yhat, average = 'weighted'), 2)
10
11 # Support Vector Machine
12 svm_yhat = best_svm.predict(x_loan_test)
13 f3 = round(f1_score(y_loan_test, svm_yhat, average = 'weighted'), 2)
14
15 # Logistic Regression
16 lr_yhat = best_lr_model.predict(x_loan_test)
17 f4 = round(f1_score(y_loan_test, lr_yhat, average = 'weighted'), 2)
18
19 f1_list = [f1, f2, f3, f4]
20 f1_list
21
```

Out[69]:

```
[0.63, 0.63, 0.76, 0.66]
```

In [70]:

```
1 # Log Loss
2
3 # Logistic Regression
4 lr_prob = best_lr_model.predict_proba(x_loan_test)
5 ll_list = ['NA', 'NA', 'NA', round(log_loss(y_loan_test, lr_prob), 2)]
6 ll_list
```

Out[70]:

```
['NA', 'NA', 'NA', 0.57]
```

In [71]:

```
1 columns = ['KNN', 'Decision Tree', 'SVM', 'Logistic Regression']
2 index = [ 'F1-score', 'Logloss']
3
4 accuracy_df = pd.DataFrame([ ll_list], index = index, columns = columns)
5 accuracy_df1 = accuracy_df.transpose()
6 accuracy_df1.columns.name = 'Algorithm'
7 accuracy_df1
```

Out[71]:

Algorithm	F1-score	Logloss
KNN	NA	NA
Decision Tree	NA	NA
SVM	NA	NA
Logistic Regression	0.57	0.57

Report

You should be able to report the accuracy of the built model using different evaluation metrics:

In []:

```
1
```